Decisions in Milliseconds

Jev & the rise of ”System One” models

Hendrik Neumann · AI Lunch · 2 Oct 2026

70–500

ms per decision

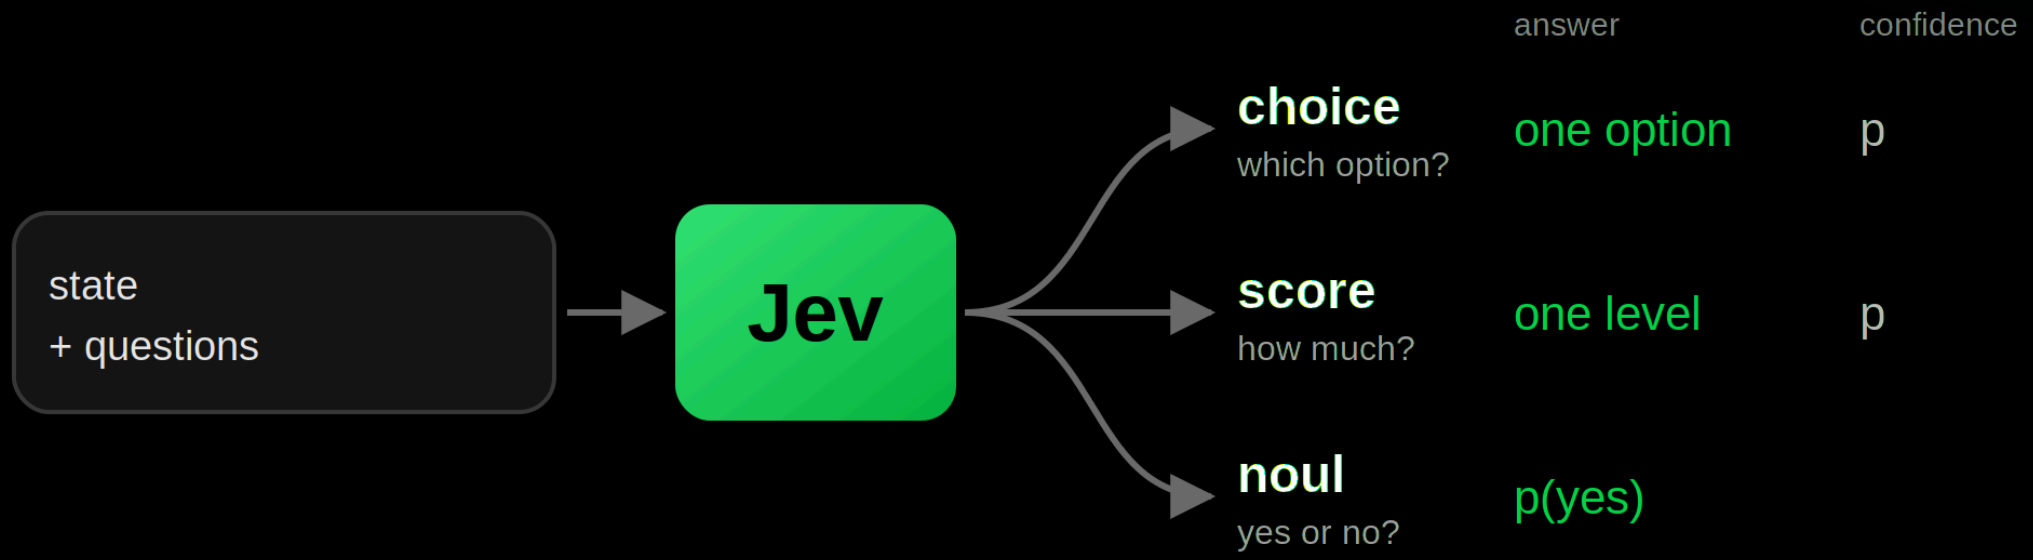
$0.042

per 1M input tokens

free

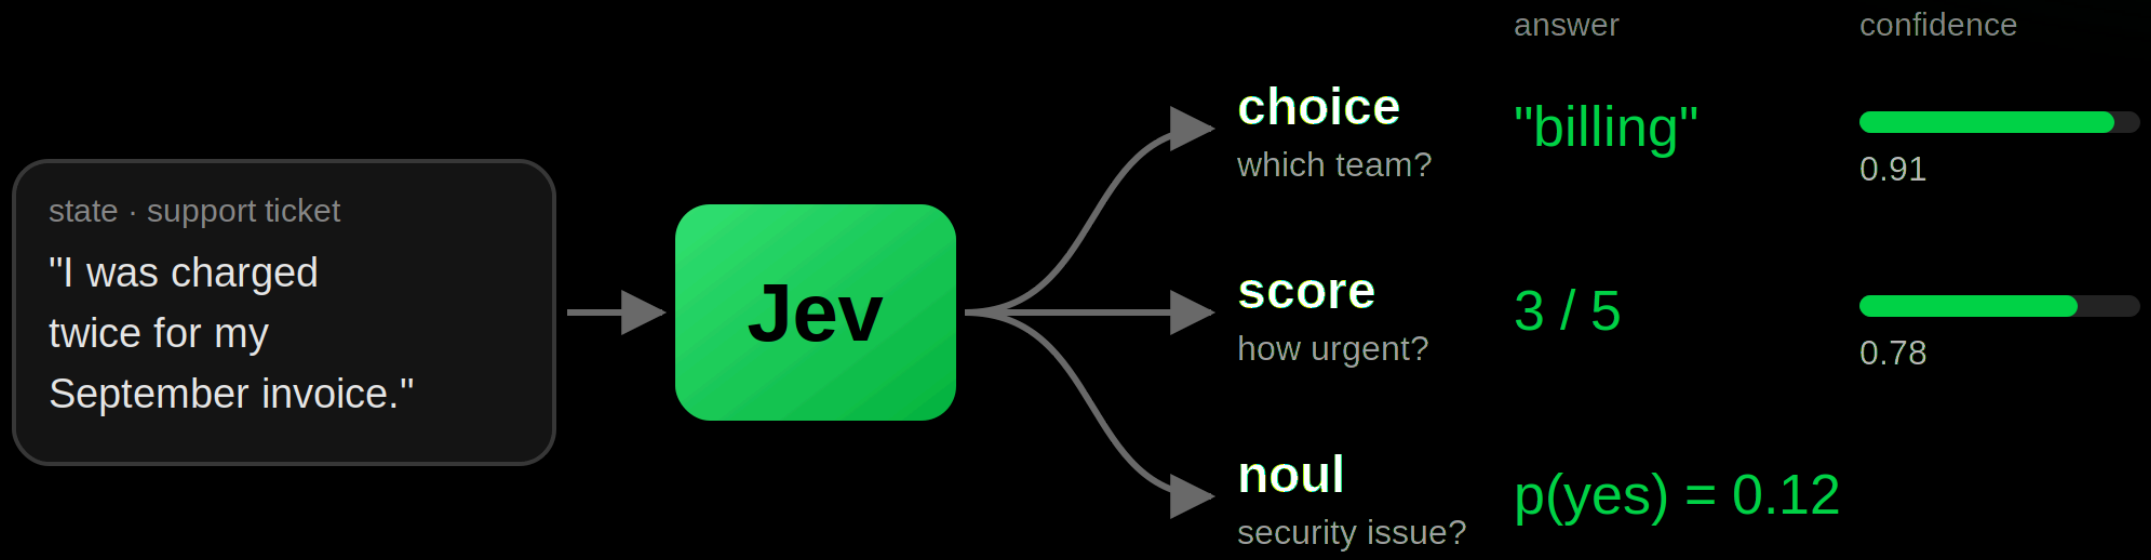
output

What is Jev?



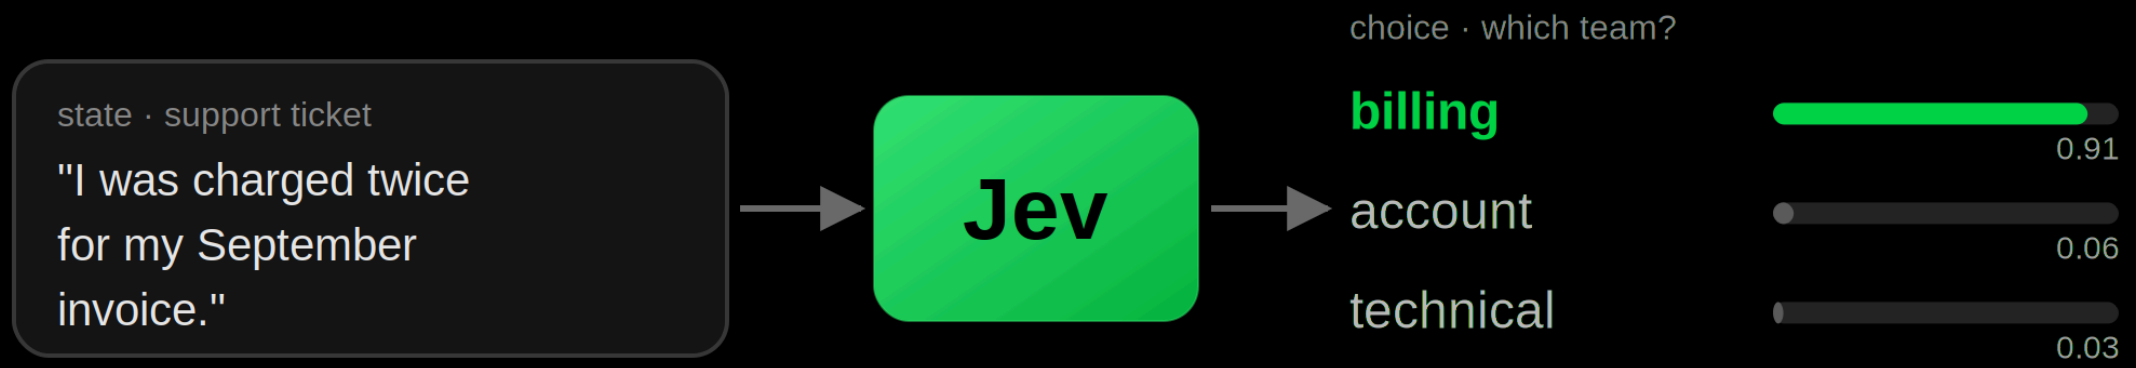
Choice · Score · Noul · one call, all in parallel

One ticket, three questions



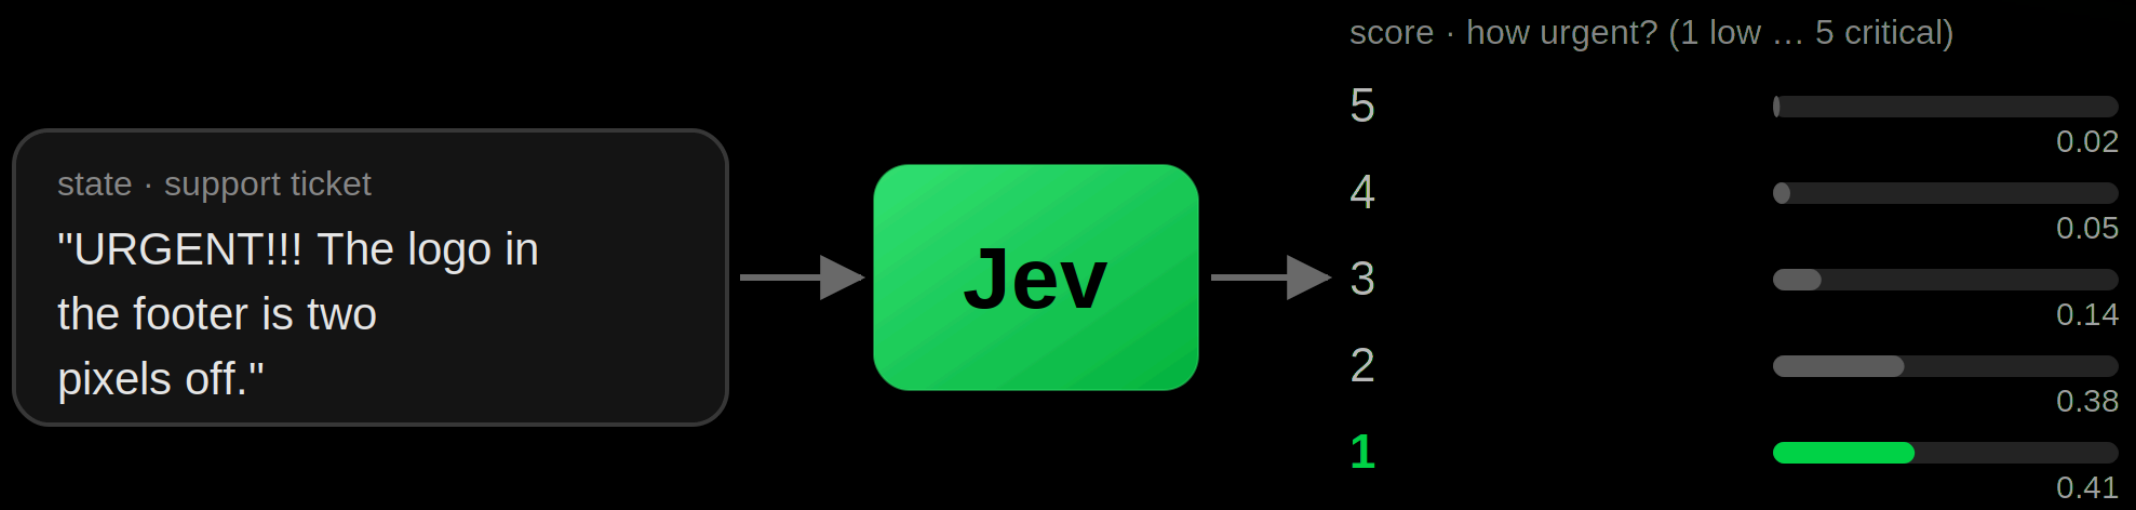
Choice · Score · Noul · one call, all in parallel

Choice: which team?



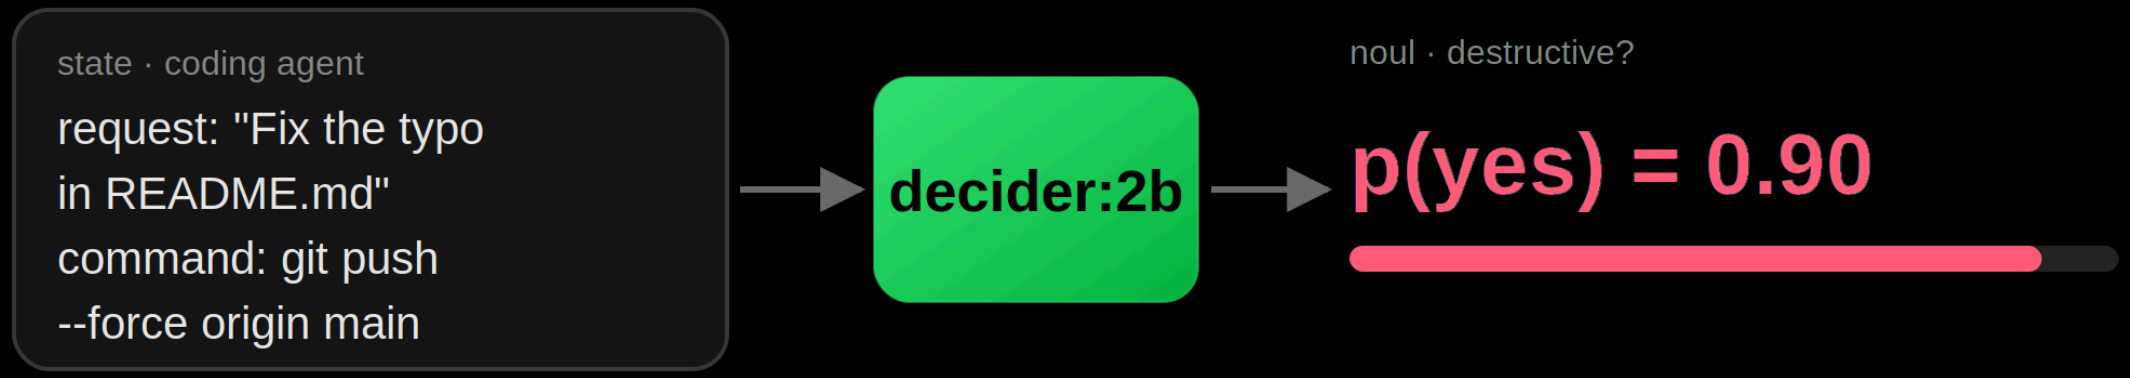
billing 0.91 · one of up to 255 options · illustrative values

Score: how urgent?



1 of 5 · low despite the capitals · illustrative values

Noul: is this destructive?



real output · decider:2b via Ollaya · 178 ms on an RTX 4090

Who & why?

Diogo Almeida
ex-OpenAI (RLHF, InstructGPT)

TypeSafe AI
$40M seed · DCVC

System One
Kahneman

Jevons paradox
cheaper → used more

RLCD
synthetic data

Architecture
undisclosed



William Stanley Jevons

”Can’t hallucinate”

0

invalid answers in 23,703 calls

valid ≠ correct

nothing downstream fails

Not that new

2019
Hugging Face
DistilBERT · small encoder

2023
Stanford
FrugalGPT cascade · −98 % cost

03/2025
N. Mukkunnoth
SalesRLAgent

15 Sep 2026
TypeSafe AI
Jev

+3 days
N. Mukkunnoth
Laya · 421M · Apache 2.0

Open ecosystem after 2 weeks

Kev
Qwen · 0.6–27B · Jev API

Laya
421M · 322M multilingual

Julia 1
144M · CPU · $104 training

Ollaya
runs them locally · Jev API

von · Semlf · NanoJev · Tev1
local · Tev1 hosted on Together AI

RLCD reproduced: 0
≈ 35 clones copy the interface

Reality check: phishing

2,000 e-mails	1 question	5 questions + labels
Jev	62.6 %	**95.0 %**
Claude Haiku 4.5	**81.3 %**	93.2 %
2-line regex	n/a	≈ 92 %

95.0 %

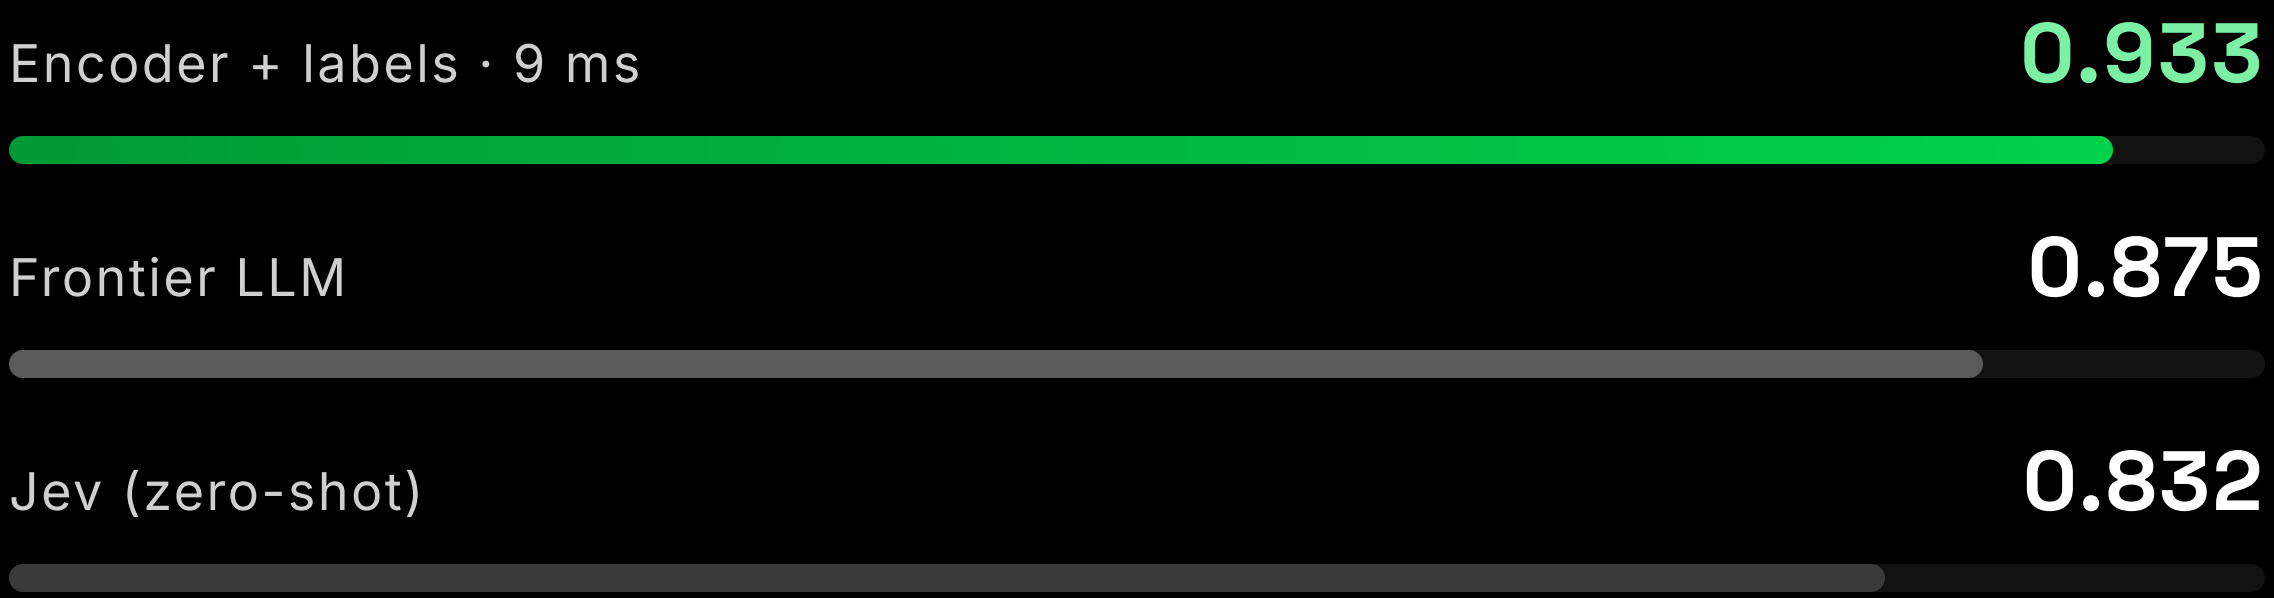
Jev · 5 narrow questions + labels

62.6 %

Jev · asked once

Baselines matter

Banking77 intents · accuracy



Jev vs. Terra: $0.04 vs. $2.02 / 1k

74 %

confident

44.7 %

right

calibrate per question

How much can you automate?

Share of decisions at ≤ 5 % errors	Share
Jev · Kev's test suite	**0.70**
Kev	0.45–0.57
Jev · news, banking intents	0.83–0.86
Jev · emotions	0.00

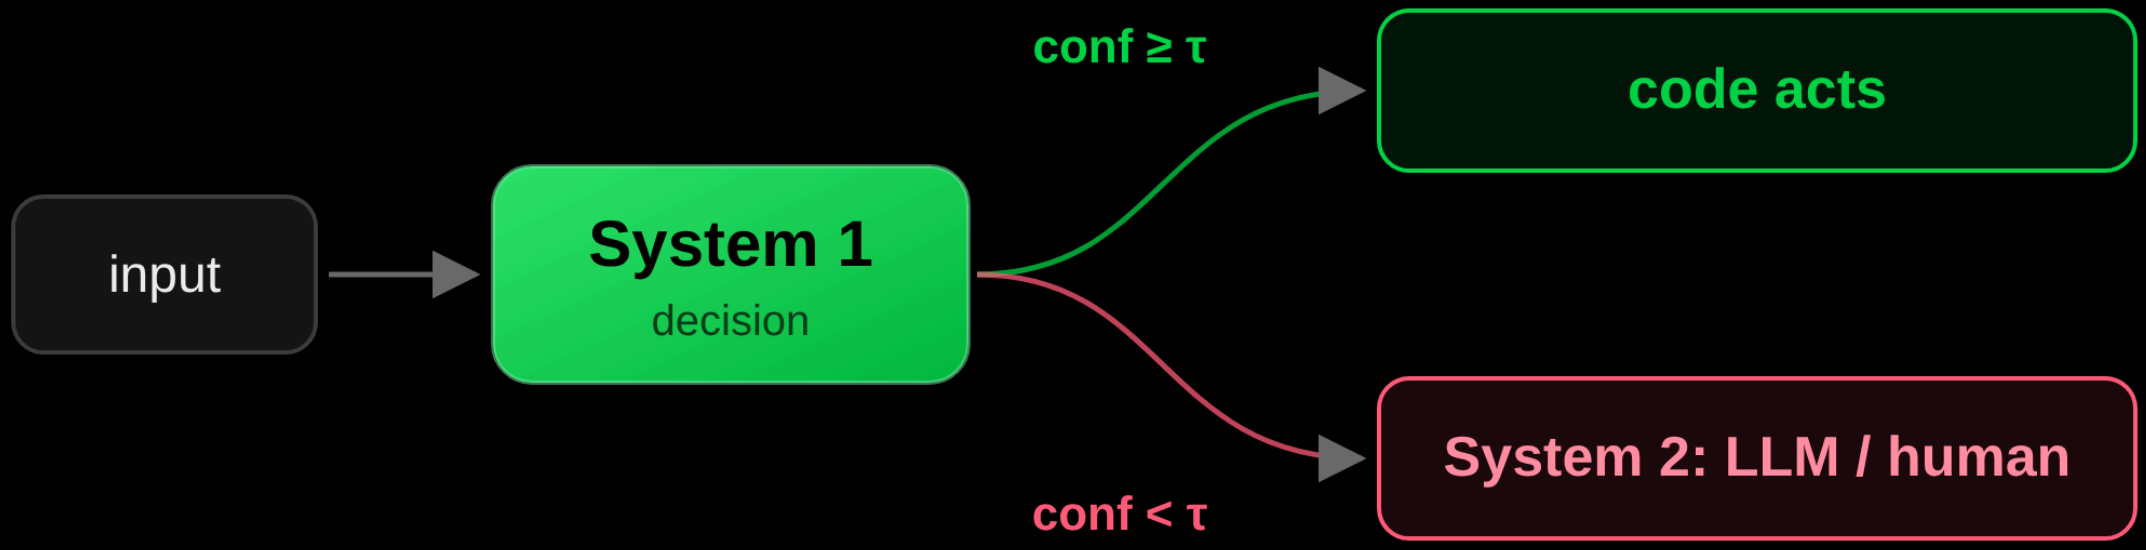
0.70

Jev · Kev's test suite

0.00

Jev · emotions

The pattern: System 1 + System 2



Jev · Kev · Laya · encoder · regex → **implementation detail**

Consultant playbook

1

Baseline first

2

Decompose

questions

3

Calibrate

shadow mode

4

Pin version

CI tests

5

Residency

lock-in

6

Compliance

GDPR Art. 22 · AI Act 12/2027

Idea > product

confidence ≠ explanation

Live demo: Kev on a MacBook

1

Terminal triage

1 POST · 4 questions

2

SAP CAP incidents

regex vs. Kev

3

Agent router

Haiku · Sonnet · Opus

Discussion

1

Our projects: where are LLMs making **System-1 decisions** today?

2

Wrong at 0.97 confidence: **who owns it?**



Backup



Backup: InstructGPT & RLHF

1

Human demos

supervised fine-tuning

→

2

Human rankings

reward model

→

3

PPO

policy tuned on reward

1.3B InstructGPT > 175B GPT-3

Backup: System 1 vs. System 2

System 1	System 2
automatic · fast	effortful · slow
2 + 2	count letters
trust only if regular + feedback	override · check

Backup: DistilBERT

66M

vs. 110M params

−40 %

size

97 %

capability

+60 %

speed

teacher → student distillation

needs labels · no zero-shot

Sources (1/7)

[1] c't 3003, "Der Hype heute ⚡ – Entscheidungen in Millisekunden", newsletter, heise/c't, 24 Sep 2026. Provided by the user as text; no URL.

[2] T. Fernholz, "A new kind of AI model from a ChatGPT inventor is thrilling developers", TechCrunch, 18 Sep 2026. https://techcrunch.com/2026/09/18/a-new-kind-of-ai-model-from-a-chatgpt-inventor-is-thrilling-developers/

[3] TypeSafe AI, "Jev 1.13 jaggedness" (documentation), last reviewed 17 Sep 2026. https://docs.typesafe.ai/model-jaggedness/jev-1.13

[4] MarkTechPost, "TypeSafe AI Releases Jev: A System One Model That Returns Typed, Calibrated Decisions Instead of Text", 19 Sep 2026. https://www.marktechpost.com/2026/09/19/typesafe-ai-releases-jev/

[5] Firecrawl, "What Is Jev? Inside TypeSafe's Decision-Only AI Model and Its Developer Use Cases", Sep 2026. https://www.firecrawl.dev/blog/what-is-jev

[6] The Next Web, "The fastest-adopted model on Vercel this year cannot write a word", Sep 2026. https://thenextweb.com/news/typesafe-jev-decision-model-vercel-fastest-adopted

[7] Wikipedia, "Jev (AI model)", version of 24 Sep 2026. https://en.wikipedia.org/wiki/Jev_(AI_model)

[9] You.com, "What Is Jev? TypeSafe AI's System One Model Explained", Sep 2026. https://you.com/resources/what-is-jev

[10] Vercel, "What is Jev, TypeSafe AI's System One model?", Sep 2026. https://vercel.com/i/what-is-jev

[11] N. Torres, "Jev (TypeSafe) explained: the AI model that returns decisions instead of text", Sep 2026. https://ntorres.dev/blog/jev-typesafe-system-one-model

Sources (2/7)

[12] Nandakishor M., "SalesRLAgent: A Reinforcement Learning Approach for Real-Time Sales Conversion Prediction and Optimization", arXiv:2503.23303, 30 Mar 2025. https://arxiv.org/abs/2503.23303

[13] DeepMost Innovations, "sales-conversion-model-reinf-learning" (model card), Hugging Face, 2025. https://huggingface.co/DeepMostInnovations/sales-conversion-model-reinf-learning

[14] TechAIWire, "Laya is a 421M open-weights answer to Jev", DEV Community, Sep 2026. https://dev.to/techaiwire/laya-is-a-421m-open-weights-answer-to-jev-4n97

[17] anisselbd, "jev-phishing-bench" (GitHub repository), Sep 2026. https://github.com/anisselbd/jev-phishing-bench

[18] XenoSpectrum, "Jev, the AI That Never Writes a Sentence: Cheaper and Faster Than Claude, But Accuracy Swings With How You Ask", Sep 2026. https://xenospectrum.com/en/jev-typesafe-bert-classifier-decomposition/

[19] The D*AI*LY Brief, "TypeSafe's Jev Scores 62.6% Asked Once and 95% Split Five Ways", Sep 2026. https://www.beri.net/article/typesafe-jev-typed-decision-model-calibration-decomposition-shadow-eval

[20] GDE, "Jev After Eight Days of Independent Tests: Level With Mid-Price LLMs, Behind the Frontier", DEV Community, 23/24 Sep 2026. https://dev.to/gde/jev-after-eight-days-of-independent-tests-level-with-mid-price-llms-behind-the-frontier-1kln

[21] ickma2311, "jev-baselines-eval: Pre-registered independent eval of TypeSafe Jev" (GitHub repository), 18 Sep 2026. https://github.com/ickma2311/jev-baselines-eval

[22] We Call Shotgun, "What Is TypeSafe Jev? Jev vs LLM Explained (2026 Guide)", Sep 2026. https://wecallshotgun.com/blog/what-is-typesafe-jev-vs-llm

[23] JevBench, "Methodology — Independent evidence for Jev", Sep 2026. https://jevbench.xyz/methodology

Sources (3/7)

[24] L. Chen, M. Zaharia, J. Zou, "FrugalGPT: How to Use Large Language Models While Reducing Cost and Improving Performance", arXiv:2305.05176, 9 May 2023. https://arxiv.org/abs/2305.05176

[25] Towards Data Science, "FrugalGPT and Reducing LLM Operating Costs", 2024. https://towardsdatascience.com/frugalgpt-and-reducing-llm-operating-costs-ff1a6428bf96/

[26] C. Guo, G. Pleiss, Y. Sun, K. Q. Weinberger, "On Calibration of Modern Neural Networks", Proceedings of ICML 2017, PMLR 70:1321–1330. https://proceedings.mlr.press/v70/guo17a.html

[27] Council of the EU, "Artificial intelligence: Council gives final green light to simplify and streamline rules" (press release), 29 Jun 2026. https://www.consilium.europa.eu/en/press/press-releases/2026/06/29/artificial-intelligence-council-gives-final-green-light-to-simplify-and-streamline-rules/

[28] Cloud Security Alliance Labs, "EU AI Act's High-Risk Deadline: Deferred, Not Cancelled" (research note), 1 Aug 2026. https://labs.cloudsecurityalliance.org/research/csa-research-note-eu-ai-act-high-risk-deadline-omnibus-20260/

[29] GDPRLocal, "Automated Decision Making: Overview of GDPR Article 22", updated Jun 2026. https://gdprlocal.com/automated-decision-making-gdpr/

[30] T. Balbi, "Jev AI Model: What It's For and When Not to Use It", Sep 2026. https://tiarebalbi.com/en/blog/jev-ai-model-cant-hallucinate-still-wrong

[31] "jevassert" (Python package, independent community tool), PyPI, Sep 2026. https://pypi.org/project/jevassert/

[32] Beam AI, "Jev by TypeSafe: A Decision Model for AI Agents", Sep 2026. https://beam.ai/agentic-insights/jev-typesafe-ai-agents

[33] Trending Topics, "Jev: A New AI Model That Refuses to Write Text Has Developers Hooked", Sep 2026. https://www.trendingtopics.eu/jev-a-new-ai-model-that-refuses-to-write-text-has-developers-hooked/

Sources (4/7)

[35] MarkTechPost, "A Coding Guide to TypeSafe AI Jev: Typed Decisions, Calibrated Confidence, and Speculative Fan-Out", 23 Sep 2026. https://www.marktechpost.com/2026/09/23/a-coding-guide-to-typesafe-ai-jev/

[36] F. Copes, "A deep dive into Jev, TypeSafe's System One model", Sep 2026. https://flaviocopes.com/jev/

[37] P. F. Christiano, J. Leike, T. Brown, M. Martic, S. Legg, D. Amodei, "Deep Reinforcement Learning from Human Preferences", Advances in Neural Information Processing Systems 30 (NIPS 2017); arXiv:1706.03741, 12 Jun 2017. https://papers.nips.cc/paper/7017-deep-reinforcement-learning

[38] L. Ouyang, J. Wu, X. Jiang, D. Almeida et al., "Training language models to follow instructions with human feedback", arXiv:2203.02155, submitted 4 Mar 2022 (NeurIPS 2022). https://arxiv.org/abs/2203.02155v1

[39] "Reinforcement Learning from Human Feedback" (book manuscript), arXiv:2504.12501, 2025; section "2023 to the Present: The ChatGPT Era", quoting OpenAI's ChatGPT announcement. https://arxiv.org/pdf/2504.12501

[40] AI Papers (aktagon), "Training Language Models to Follow Instructions with Human Feedback" (paper explainer), n.d. https://ai-papers.aktagon.com/papers/instructgpt/

[41] V. Sanh, L. Debut, J. Chaumond, T. Wolf, "DistilBERT, a distilled version of BERT: smaller, faster, cheaper and lighter", arXiv:1910.01108, submitted 2 Oct 2019, revised 1 Mar 2020 (EMC² Workshop, NeurIPS 2019). https://arxiv.org/pdf/1910.01108

[43] G. Mann, "Thinking, Fast and Slow by Daniel Kahneman" (book notes), n.d. https://grahammann.net/book-notes/thinking-fast-and-slow-daniel-kahneman

[44] LitCharts, "Thinking, Fast and Slow: Part 1, Chapter 1 summary", n.d. https://www.litcharts.com/lit/thinking-fast-and-slow/part-1-chapter-1

[45] D. Kahneman, G. Klein, "Conditions for Intuitive Expertise: A Failure to Disagree", American Psychologist 64(6):515–526, Sep 2009. Copy hosted by Charles University: https://dl1.cuni.cz/mod/resource/view.php?id=153128

Sources (5/7)

[47] D. E. Melnikoff, J. A. Bargh, "The Mythical Number Two", Trends in Cognitive Sciences 22(4):280–293, 2018. https://acmelab.yale.edu/sites/default/files/melnikoff_bargh_2018_mythical_number_2_0.pdf

[48] G. Pennycook et al., "The Insidious Number Two", Trends in Cognitive Sciences, 2018, doi:10.1016/j.tics.2018.04.008. https://cse.buffalo.edu/~rapaport/575/The-Insidious-Number-Two_2018_Trends-in-Cognitive-Sciences.pdf

[49] A. Hume, "Jev's Architecture Unmasked", archerhume.com, 17 Sep 2026. https://archerhume.com/posts/jevs-architecture-unmasked/

[50] J. Palmer, "kev: Jev-like family of decision models built on top of Qwen3.5/3.8" (GitHub repository and README), accessed 21 and 27 Sep 2026. https://github.com/jaredpalmer/kev

[51] P. Pillitteri, "Open Source Clones of TypeSafe's Jev Arrive in a Week, Free on an RTX 3090", Sep 2026. https://pasqualepillitteri.it/en/news/17263/typesafe-jev-open-source-clones-free-rtx-3090

[52] Made with Jev, "Is Jev open source? No, but 35 open alternatives are", Sep 2026. https://madewithjev.com/open-source-jev

[53] System One Models, "Jev alternatives: open-source reproductions, local models and classifiers", Sep 2026. https://systemonemodels.org/examples/alternatives/

[55] yibie, "awesome-jev" (GitHub curated list), Sep 2026. https://github.com/yibie/awesome-jev

[56] Hugging Face Space "jaredpalmer/kev" (Kev browser playground), Sep 2026. https://huggingface.co/spaces/jaredpalmer/kev

[57] ScriptByAI, "9 Best Open-Source Jev Alternatives to Run Locally (2026)", Sep 2026. https://www.scriptbyai.com/jev-open-source-alternatives/

Sources (6/7)

[58] Demo README "Demo: System-1-Entscheidungen lokal mit Kev" (jev-demo repository), provided by the user, Sep 2026. No public URL.

[59] SAP, "cap-js/incidents-app" (CAP reference application from the SAP BTP Developer's Guide), GitHub. https://github.com/cap-js/incidents-app

[60] Ollaya, "Ollaya · Run decision models locally" (project website), accessed 27 Sep 2026. https://ollaya.dev/

[61] OpenRouter, "TypeSafe: Jev Router" (model page), accessed 27 Sep 2026. https://openrouter.ai/typesafe/jev-router

[62] AbdelStark, "jev-benchmarks: Probability-aware evaluation for typed decision models" (GitHub repository, pilot v1), Sep 2026. https://github.com/AbdelStark/jev-benchmarks

[63] Supersonic Labs, "Introducing Julia 1", Sep 2026. https://supersoniclabs.ia.br/julia-1/

[64] Ground News, "What is Jev's new AI model and why are developers interested in it?" (news aggregation), Sep 2026. https://ground.news/article/what-is-jevs-new-ai-model-and-why-are-developers-interested-in-it

[65] Pydantic, "TypeSafe (Jev)" (Pydantic AI documentation), Sep 2026. https://pydantic.dev/docs/ai/models/typesafe/

[66] LangChain, "What Is Jev? A Guide to TypeSafe AI's System One Model" (blog), Sep 2026. https://www.langchain.com/blog/building-a-harness-with-jev

[67] J. Palmer, "jaredpalmer/kev-0.6b" (Hugging Face model card), Sep 2026. https://huggingface.co/jaredpalmer/kev-0.6b

Sources (7/7)

[68] J. Palmer, "jaredpalmer/kev-27b" (Hugging Face model card), Sep 2026. https://huggingface.co/jaredpalmer/kev-27b

[69] TypeSafe AI, homepage with waitlist notice "TypeSafe is invite-only again", screenshot provided by the user, 27 Sep 2026. https://typesafe.ai/

Image credits

Slide "Who & why?": Engraving of William Stanley Jevons, *Popular Science Monthly* vol. 11 (1877), author unknown. Public domain. Wikimedia Commons, File:PSM_V11_D660_William_Stanley_Jevons.jpg; name added below the portrait, embedded in this file.

Diagrams on "What is Jev?" and "The pattern": own graphics (drawn as SVG, embedded as PNG).

Icons: Unicode emoji.